

Security & Privacy

Stephen Farrell

`stephen.farrell@cs.tcd.ie`

Course materials:

<https://down.dsg.cs.tcd.ie/cs7053/>

<https://github.com/sftcd/cs7053>

Slideware + some papers

Computer Security Concepts

*Introduction of some basic
concepts and terminology in
computer security*

Next hour(s)...

- A bit more history...
- Application layer security
- Security evaluation
- Network security
- Identification & Authentication

Computer Security Goals

Traditional Security Goals: "CIA"

- Confidentiality
 - Keeping secrets
- Integrity
 - Preventing unauthorised modifications
 - Keeping data consistent
- Authentication/Assurance/Availability
 - The meaning of the “A” depends upon who you ask!

Services and mechanisms

- From the old “OSI” world, we've inherited two (sometimes confusing) concepts:
 - Security service: provides a security function to the system, based on the use of security mechanisms, e.g. confidentiality
 - Security mechanism is a technique, or protocol etc. which can be used to provide a service, e.g. encryption, access control

Services/Mechanisms

- Confidentiality
 - Access control
 - Cryptography
 - Secure communication channels
- Integrity
 - Checksums
 - Digital signatures
 - Secure communication channels
- Availability
 - Designing secure software and communication protocols
 - Load balancing
 - Anti-DDoS services

Structure of Military Security

- Information is classified according to national security
- The classification is clearly labeled on the binder
- All classified information is stored in a safe
- All users are “cleared” to see information up to a certain level
- Users have to prove their clearance to withdraw the binder from the safe
- Additional compartments enforce the *need-to-know* principle

Creating New Information

- New files are labelled with the classification of the most secret component
- Aggregation of unclassified information may generate a “top secret” file
- Sanitization downgrades the label of existing information

Old Problems Aggravated

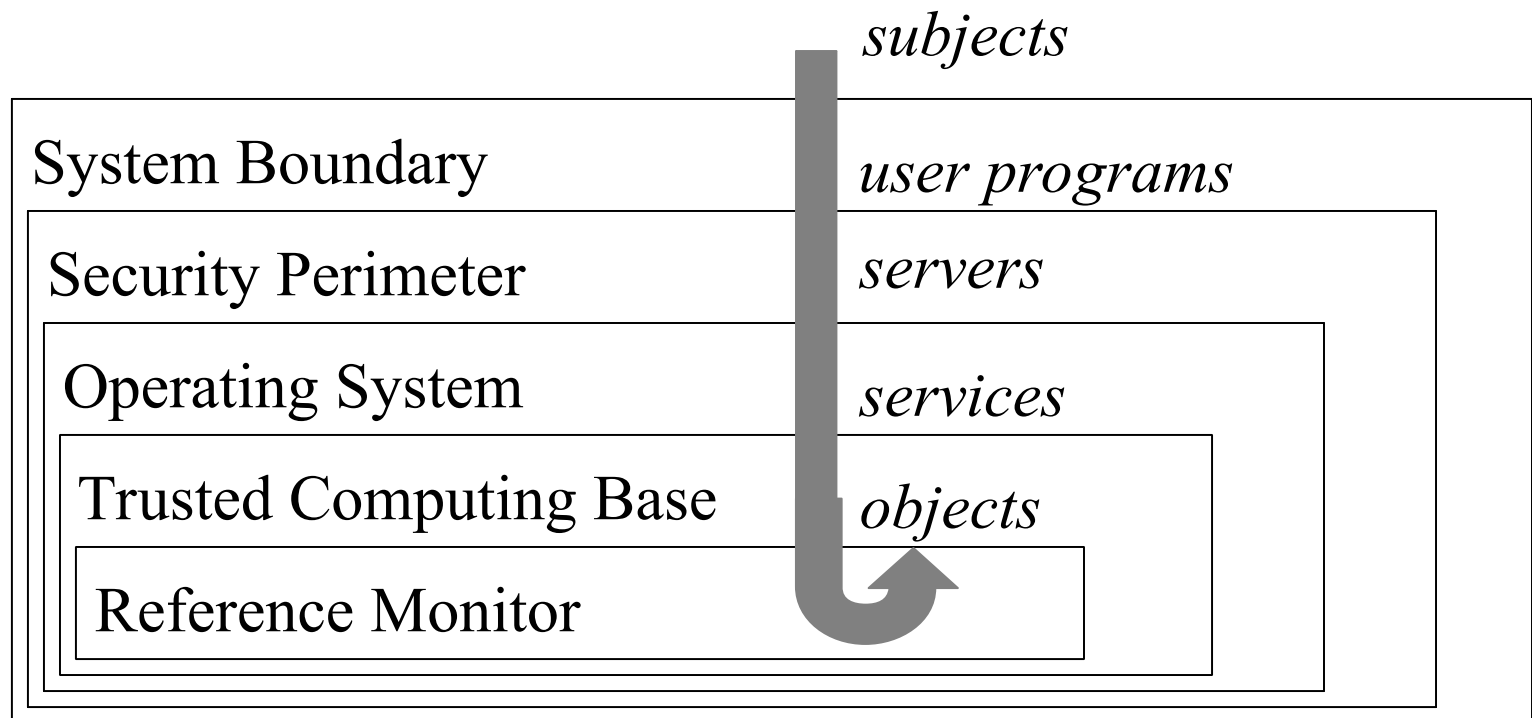
- Aggregation
 - it is easier to aggregate information from a vast set of information
- Authentication
 - it is more difficult for a computer to identify a person (partly solved)
- Browsing
 - it is easier to read all files in a file system than all files in a safe
- Integrity
 - modification is harder to detect (mostly solved)
- Copying
 - digital copies are indistinguishable from the originals (really there are no originals at all!)
- Denial of Service
 - denial of service is a notorious characteristic of computer systems

Covert Channels

- Hidden means of communication, that allows information to be leaked to third parties
- Some types:
 - Timing Channels
 - observable differences in system utilization
 - Inference Channels
 - intersection of non-classified information
 - Fabrication Channels
 - aggregation of non-classified information

Security Perimeters

users, remote workstations, the Internet



Attack Surface

- Intuition: more possible ways in which attacks can be attempted => more likely an attack may succeed
 - Listener on a port, Callable API, Command line arguments,...
- Think of the potential attack points as if they formed a geometric construct: the attack surface
 - Smaller attack surface == better! (probably)
 - Useful relative metric as you evolve a system
 - Argues (to an extent) against over-generic tools/re-use (e.g. CMS for web sites perhaps)
- Dec 2018 survey paper (paywall):
 - <https://doi.org/10.1016/j.infsof.2018.07.008>
- OWASP guidance:
 - https://www.owasp.org/index.php/Attack_Surface_Analysis_Cheat_Sheet
- Michael Howard concept from 2003:
 - [https://docs.microsoft.com/en-us/previous-versions/ms972812\(v=msdn.10\)](https://docs.microsoft.com/en-us/previous-versions/ms972812(v=msdn.10))

Sidebar: Operating Systems

- Won't cover OSES or virtualisation much as issues are very system dependent
 - UNIXes, incl. GNU/Linux, OpenWRT, Apple iOS/OS X, Android, flavours of busybox, Windows, Cisco IOS, JunOS,...
- Basic idea is to isolate processes (users are represented as processes) and privileges
- Cross-boundary vulns have most impact
- Local privilege escalation attacks: meh
- Remote code execution (RCE) or break out of VM/browser sandbox: BIG deal!
- Recent work on side-channels (spectre/meltdown etc.) tells us speed and security trade-offs have been mismanaged in many environments for many years

Commercial Security

- Historically the military paid more for security, were more accepting of the inconvenience (which almost all security mechanisms introduce), and had a chain-of-command upon which to fall back
- Co-operating/competing commercial enterprises have none of these

The 1988 Internet worm

- Propagated (mainly) via SMTP
 - using the DEBUG build of sendmail (then the prevalent mail server)
 - Also exploited fingerd problems
- It killed the Internet for a whole day!
- Spafford: “The Internet worm program: An analysis” well worth a read

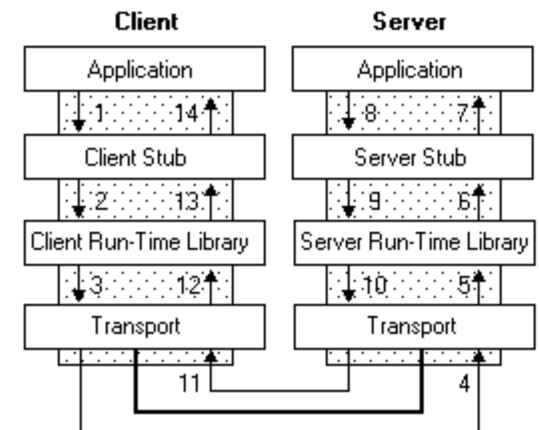
<https://down.dsg.cs.tcd.ie/cs7053/materials/spafford88internet.pdf>

Typical Enterprise Security Model (circa 1995)

- System and network administrators setup and manage users and applications
- DCE exemplifies this approach
 - <http://www.opengroup.org/dce/>
 - Interesting that URL still works
 - But code download seems to have disappeared
 - Part(s) maybe survive at <https://github.com/dcerpc/dcerpc> or <http://opendce.hands.com/>
 - Kerberos (RFC4120) – still persists as part of MSFT windows and a few other systems

DCE RPC Model

- Server has function $f(a,b)$; client wants to call
- Interface definition language (IDL) used to generate client and server “stubs”
 $f_client(a,b)$ and $f_server(a,b)$
- Run-time hides transport from client
- Client call to $f(a,b)$ goes around the world
- Transport does security stuff (via GSS-API in the case of DCE RPC)
- Diagram from:
<https://docs.microsoft.com/en-us/windows/win32/rpc/how-rpc-works>
- Diagram doesn't show Kerberos server(s)



Problems with the 1990's enterprise security model

- Generally assumed a homogeneous network and set of applications
 - Became more false as time went on
- Assumption that all users and applications were centrally managed
- Deployment showed up performance and usability issues

Security APIs

- Periodically someone tries to develop an API to “hide” security
- Has worked fine for cryptographic primitives
 - MS-CAPI, PKCS#11, JCE, NaCL, ...
 - Pretty much all good development tools and environments have one (python, PHP, ...)
- Not so successful for higher layer functions
 - GSS-API, SPKM

Sidebar: Crypto makes interop hard

- It is **much** easier to get a system to (supposedly) “work” in “insecure” mode
 - And **much, much** easier when secure mode involves cryptography
- This is a general problem which has lead many people to turn off/down security
- Answer: develop with crypto turned on
- Bear it in mind as you design things
- Make sure you’ve done that at least once!

Along came the web

- The number of connected hosts rose exponentially for a while
- Highlights security issues with:
 - Proxies (various bad things can happen at a proxy)
 - Tunnelling (“protocol-X” over HTTP)
 - Having a ubiquitous tunnel end-point on many machines
 - Browser security models

Security Evaluation and Network security

*Customers need some
confidence that the system or
network that they are about to
purchase is “secure”.*

Security Evaluation Criteria

[TCSEC] Department of Defence Trusted Computer Systems
Evaluation Criteria (Orange Book)

- [https://csrc.nist.gov/publications/secpubs/rainbow/OOPS! Jan 2019 US Govt. Shutdown ;-\)](https://csrc.nist.gov/publications/secpubs/rainbow/OOPS!%20Jan%202019%20US%20Govt.%20Shutdown%20;-))

[ITSEC] IT Security (UK, NL, FR, DE)

https://www.bsi.bund.de/SharedDocs/Downloads/DE/BSI/Zertifizierung/ITSicherheitskriterien/its-ec-en_pdf.pdf?__blob=publicationFile

[Others] CTCPEC (Canada) JCSEC (Japan)**[CC]**

[CC] Common Criteria (ISO IS 15408)

- <https://www.commoncriteriaportal.org/files/ccfiles/CCPART1V3.1R5.pdf>

Security Evaluation

- How do we compare the security of different computer systems?
 - Different authentication mechanisms
 - Kerberos, passwords, smart cards, ...
 - Different access control mechanisms
 - Access Control Lists, Role-Based Access Control, ...
 - Different cryptographic algorithms
 - DES, AES, RSA, ECC, D-H, Curve25519...

“Orange Book” Classification

Class	Title	Key Features
A1	Verified Design	Formal top-level specification and verification, formal covert channel analysis, informal code correspondence demonstration
B3	Security Domains	Reference monitor (security kernel), “highly resistant to penetration
B2	Structured Protection	Formal model, covert channels constrained, security oriented architecture, “relatively resistant to penetration”
B1	Labeled Security Protection	Mandatory access controls, security labeling, removal of security related flaws
C2	Controlled Access Protection	Individual accountability (authentication) and extensive auditing
C1	Discretionary Security Protection	Discretionary access controls, protection against accidents among cooperating users
D	Minimal Protection	Not Classified

ITSEC

- Separation of functionality and assurance
- Functionality-classes (F-C1 – F-B3)
corresponds to the Orange Book
- Assurance-classes (E0 – E6)
 - E0: lowest, E6 highest
 - E4: commonest (maybe)
- Target of Evaluation
 - Set of evaluated components: security policy, security related functions, definition of the required security mechanisms for the target level of evaluation

Common Criteria

- Combines previous criteria
 - Separation of functionality and assurance
 - Target of Evaluation (TOE)
 - Security Target (ST) = desired level of evaluation
- Protection Profiles (PP)
 - Security Requirements
 - Security objectives
 - Independent of Implementation

Other aspects of assurance

- Installation and use must be considered
 - What are the defined security policies?
 - How is the software installed?
 - How are the machines administered?
- Development tools & environments
 - SNMP exploit due to buggy 3rd party ASN.1 handling library

Assurance != "Works-as-advertised"

- Assurance only shows that the product/system matches a specification
 - NOT that it does what the customer wants!
- The product owner pays the evaluator (lots!!)
 - So although the evaluator is Government licensed, they have reasons to be nice to the developer.
- Caveat emptor rules ok.

Network security

- Different from application layer security
 - Usually no real interest in APIs here
 - Normally trying to secure either a (sub)network, or network node from network based attacks
 - Frequently network nodes don't use standard operating systems
 - E.g. Cisco routers, 802.11 access points

Network Security

- First, read about networks, you need to have some level of understanding of IP (both), DNS, TCP, BGP, HTTP and how e.g. Javascript/PHP works on the web
 - Go find your own URLs:-) Or a book.
- I expect you'll do this yourself or have gotten it from other modules from this or other years' study
- Some good network security reading starters:
 - <https://www.sans.org/> training stuff etc.
 - <https://isc.sans.edu/> "Internet Weather"
 - <https://first.org/library/> more reading
 - <https://www.owasp.org/> web security

Open source intelligence

OSINT

MISP - Malware Information Sharing Platform
and Threat Sharing

<https://www.misp-project.org/>

<https://github.com/MISP>

Idea is to share information about attacks
and indicators of compromise (IOCS)

Some events shared openly, some less so

You can browse some JSON at

<https://www.botvrij.eu/data/feed-osint/>

Learn from networks you use...

- Browser developer UI to see HTTP/JS
 - network/JS/console: “shift-ctrl-I” or similar
- Wireshark/tshark/tcpdump to see packets
 - Wireshark dissectors exist for many protocols, DNS, TLS, ...
- Protocol capture (PCAP) files
 - Careful of privacy – PCAPs can be revealing and are hard to anonymise
 - Please tell me if you know of good up-to-date PCAP anonymisation tools!
- If this is new to you - **do play with these** – you’ll learn stuff!

Network security view

- Original Internet “architecture” assumed end-to-end connectivity
- Hence end-to-end security was the main consideration for those developing the Internet
 - But they were **very** slow developing IPsec (about 10 years!)
 - And meanwhile NAT and firewalls arrived

The End-to-End Argument

Saltzer, Jerome H., David P. Reed, and David D. Clark. "End-to-end arguments in system design." ACM Transactions on Computer Systems (TOCS) 2.4 (1984): 277-288.

<https://dl.acm.org/citation.cfm?id=357402>

READ THAT!!!

And then RFCs 1958 & 3439

And RFC 7258 (why not:-)

— caution, author present

E2E Argument has current consequences...

HTTP/2 and QUIC are protocols that attempt to have e2e encryption as a mandatory mechanism (ish) not (only) in order to achieve a confidentiality service but primarily to mitigate ossification

Some, but not all, of those who know what's going on are happy with that

FWIW, (which is not much) I am happy with that.

Network Address Translation (NAT)

- NAT is today mainly used to hide local addresses from the Internet for (mostly) provisioning reasons
 - NAT means that the higher layer end-points “see” different addresses for one another
 - Breaks many end-to-end assumptions
- “Carrier Grade” NAT (CGN) is coming to an ISP near you soon (or has already)

CGN & Security

Many more endpoints behind each public IPv4 address; put a NAT box in the ISP network

NAT444

There are other, maybe better, IPv6 transition mechanisms

Some think this is an IPv6 avoidance mechanism

Good way to freak out “legal intercept” (LI) fans

LI+CGN: Good way to freak out ISPs who have to log

Good way to freak out anyone who wants to offer a public service

PITA for DNSSEC

Bad way to get to an IPv6 world?

But may be needed for some time

NAT problems with SIP

UserA	NAT	NS	UserB
10.1.1.221	192.168.221.1	63.88.221.88	192.168.1.10
----F1 INVITE---->	--F2 INVITE->		
		--F5 INVITE->	
<---F4 100-----	<-F3 100-----		
		<--F6 180----	
<---F8 180-----	<-F7 180-----		
		<--F9 200 OK-	
<---F11 200 OK----	<-F10 200 OK-		
----F12 ACK----->	---F13 ACK---		
		--F14 ACK--->	
<---F15 2WAY RTP->	<-----F16 2WAY RTP----->		
----F17 BYE----->	-F18 BYE----		
		--F19 BYE--->	
		<-F20 200 OK-	
<---F22 200 OK----	<-F21 200 OK-		

Firewalls

- Until the mid-90's most sites didn't bother filtering traffic
- It became clear that exposing your internal network topology could cause problems
 - E.g. If I know that router brand “X” has an open default configuration and I can see that you've got one of those, then my attack point is clear

Filtering routers

- Initially people put in filtering rules in border routers
 - E.g. “no packets with a destination on my inner network are allowed out”
- But, IP spoofing attacks meant that this wasn't sufficient
 - So firewall products developed
 - Note: has been biggest security products market, I've not checked recently

Denial-of-Service (DoS)

- Consume resources to make a service unavailable
 - Has been known as a vulnerability for many years
 - Began to be exploited in last decade on major Internet sites
 - DDoS attacks against the DNS roots
- DDoS = Distributed Denial of Service

TCP SYN flooding attack

- SYN packets with spoofed IP addresses cause the server to maintain state
- Flooding the server that way is your basic DoS attack
 - CERT Advisory CA-96.21

<https://mailman.apnic.net/mailling-lists/apple/archive/1996/09/msg00131.html>

Network access protocols

- Issue is how to decide when to allow a node to join/use the network
 - IEEE 802.1x networks (eduroam)
 - Corporate networks
 - Mobile operator networks
- RADIUS protocol (or Diameter)
 - RFC2865 (or RFC6733)

Search Engines

- ...to find vulnerable systems
- Shodan HQ <https://www.shodan.io/>
 - Est. 2013 - search for “camera”
- Censys <https://censys.io/>
 - Mainly IPv4 based surveying (zmap)
- Ichidan ichidanv34wrx7m7.onion
 - Didn't work for me
- There are more I guess...

Tor

- Onion routing
<https://www.torproject.org/>
- Tor Browser (TB)
<https://www.torproject.org/projects/tor-browser.html.en>
- Tor Hidden services: “<hash>.onion”
DNS-like names
- Search engine: <https://ahmia.fi/>
- Be a bit careful though!

My browser setup

- TBB: If searching for something sensitive (e.g. medical info)
- Search default: DuckDuckGo (Where “!g” works too:-)
- Default browser: FF + NoScript/Ghostery & disallowing cookies, with some white-listed sites
- Some sites don't work with the above; mostly: screw 'em
- If-need-be: chromium/incognito with a chmod of `$HOME/.config/chromium` to make almost all read-only (but not sure of effectiveness)
- If-all-else-fails: Brave or Vivaldi
- On phone: sailfish browser with no JS/no cookies and 2ndary open-kimono browsers if-need-be (Webcat/Web pirate)
- Recommend you figure out some browser-hygiene you prefer and implement that – be willing to help others do the same!

Identification and Authentication

Establishing identity and
verifying credentials

Identification

- Establish the identity of principals by means of:
 - something known (*password, PIN*)
 - something possessed (*smart card, Java ring*)
 - something personal (*fingerprint, retina scan*)
 - something to do (*signature*)
- Combinations of above (*smart card + PIN*)

Passwords

- Most systems rely exclusively on passwords

login: *username*

password: *******

- Password scheme problems:
 - users choose “bad” passwords (*password*, *secret*)
 - compromised via sniffing
- Good passwords won't be in a dictionary
- There are better password protocols
 - but first we need some more crypto

2008 Count of some of my passwords

Category	Count	Known	Examples
Logins	10	6	Laptops, host systems (incl. root accounts)
Devices	5	0	DSL router, home print/file servers, sensor nodes
Network access	4	0	Work n/w, WLAN, ISP, etc
Protocol	14	1	Outbound HTTP proxy, IMAP, Jabber, skype, etc.

Biometrics

- Many people (apparently) think that replacing username & password with a biometric is a great idea.
- It might be:
 - If secure biometrics exist
 - If the re-use implications are acceptable
 - If the context supports the full life-cycle

Biometric methods

- Fingerprint
- Retina scan
- Face recognition
- Gait (walking)
- Toe-smell

Fingerprint

- Prof. Tsutomu Matsumota's well publicised 2002 attack against most common fingerprint recognition engines
 - <https://cryptome.org/gummy.htm>
 - Information here is directly from his paper: "Impact of Artificial "Gummy" Fingers on Fingerprint Systems"

Fingerprint scanner

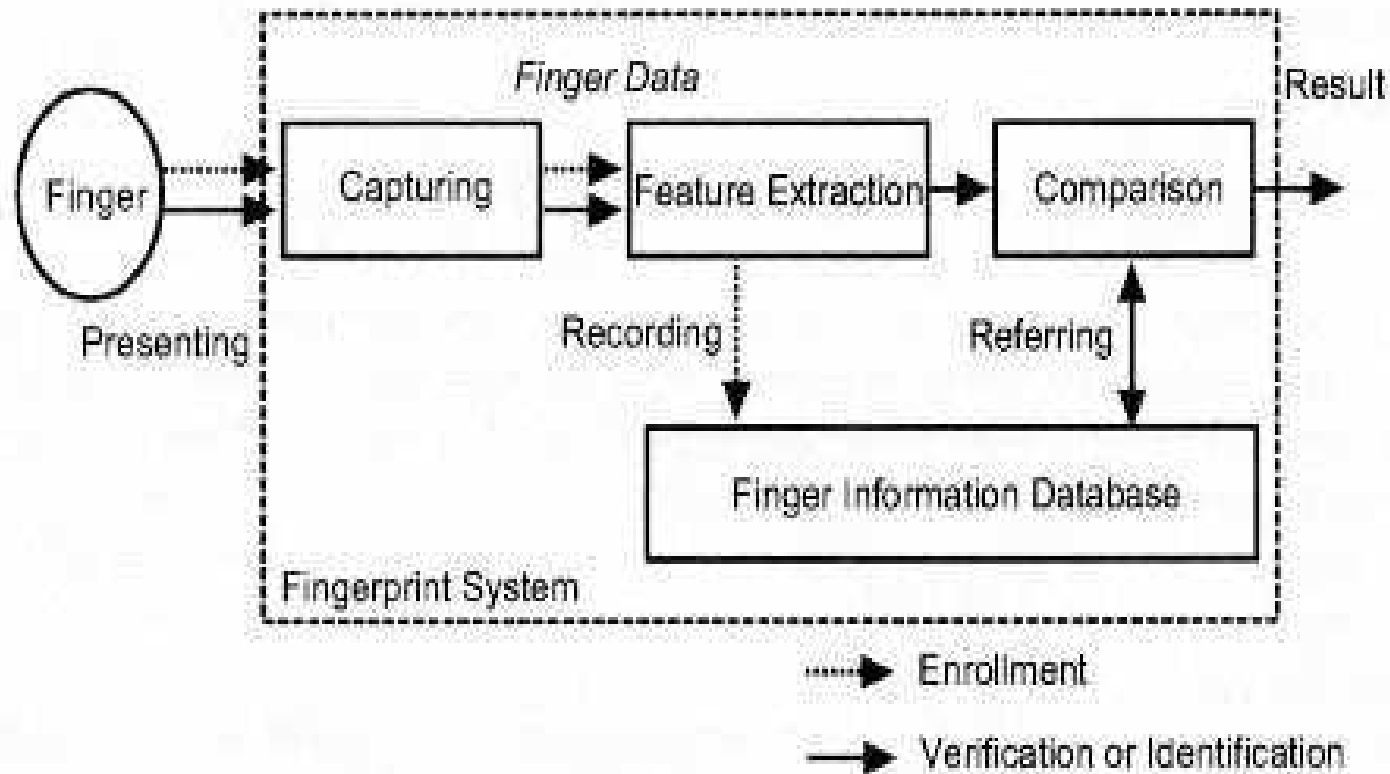




Figure 2.1 Typical structure of a fingerprint system

Make a gummy finger



(a) The mold for gummy fingers



(b) Gummy finger

Figure 4.5 Photographs of the outside appearance of the mold and a *gummy* finger. The *gummy* finger was produced from a residual fingerprint on a glass plate, enhancing it with a cyanoacrylate adhesive



Figure 4.6 The Fingerprint image of the *gummy* finger, which was displayed by the system with Device H (equipped with a capacitive sensor).




Figure 4.7 Average number of acceptance for each device, in terms of *gummy* fingers which were cloned from residual fingerprints. Here, the subject is one person.

Results

- 11 of 11 scanners tested were broken
 - The list was extended a bit more since
- That's enough of a result!

But there's more! ...

CCC 2017: 55 minute video of breaking biometrics

https://www.youtube.com/watch?annotation_id=annotation_2684251971&feature=iv&src_vid=ply6k4gvQsY&v=VVxL9ymiyAU

Fiebig, Tobias, Jan Krissler, and Ronny Hänsch.
"Security Impact of High Resolution Smartphone
Cameras." WOOT. 2014.

<https://www.usenix.org/system/files/conference/woot14/woot14-fiebig.pdf>

When are biometrics ok?

- What applications?
 - What type(s) of biometric?
 - What benefits?
 - What costs?
 - Financial and other (e.g. Privacy)

Privacy

- Rad RFC6973 - In addition to “normal” security threats we need to care about
 - Correlation
 - Identification
 - Secondary use
 - Disclosure
 - Exclusion
 - **Re-identification**
- Actually, it's not entirely clear to me that the risk analysis methodology we follow for security works well for privacy
 - There's scope for research there

Privacy Puzzle

- Emails contain a Received header field which can contain the mail user agent IP address
 - What consequences?
 - Overall good or bad from a privacy perspective?

What is your #1 mitigation for
all problems?

What is your #1 mitigation for all problems?

Yes, Backup.

Do that. Early and often.

Summary

- Services vs. Mechanisms, CIA
- ~30 year history
- Common Criteria
- N/W security a bit different from system security
- Firewalls etc.
- Passwords (yuk)
- Biometrics (also yuk)
- Privacy (yay!)
- Do backups